

Database Systems (資料庫系統)

12/5/2022

Database #2

Introduction to Computer (計算機概論)

Winston Hsu, NTU

1

Administrative Matters

- HW#8 (SQL queries) due on December 12, 2022 (12pm, Wednesday)
- Final exam on December 19, 2022
 - Range: Python programming (lecture 9) - Database (II) (Lecture 15)

2

Reflection: DB design

- Step 1: Requirements Analysis
 - What data to store in the database?
- Step 2: Conceptual Database Design
 - Come up with the design: Entity-Relation (ER) model
 - Sketch the design with ER diagrams
- Step 3: Logical Database Design
 - Implement the design: relational data model
 - Map ER diagrams to relational tables

3

Example Table Definitions

*Sailors(sid: integer, sname: string, rating: integer,
age: real)*

Boats(bid: integer, bname: string, color: string)

Reserves(sid: integer, bid: integer, day: date)

4

SQL: Queries, Constraints, Triggers

Chapter 5

5

Lecture Outline

- Basic Query
 - *SELECT*
- Set Constructs
 - *UNION, INTERSECT, EXCEPT, IN, ANY, ALL, EXISTS*
- Nested Queries
- Aggregate Operators
 - *COUNT, SUM, AVG, MAX, MIN, GROUP BY, HAVING*
- Null Values
- Integrity Constraints
 - *CHECK, CREATE ASSERTION*
- Triggers
 - *CREATE TRIGGER, FOR EACH ROW*

6

Example Table Definitions

Sailors(sid: integer, sname: string, rating: integer, age: real)

Boats(bid: integer, bname: string, color: string)

Reserves(sid: integer, bid: integer, day: date)

- Find names of sailors who've reserved boat #103

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=103
```

7

Cross-Product

- Each row of S1 is paired with each row of R1.
- Result schema** has one field per field of S1 and R1, with field names 'inherited' if possible.
 - Conflict: Both S1 and R1 have a field called sid.
referred to solely by position

S1				S1 x R1						
sid	sname	rating	age	(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	dustin	7	45.0	22	101	10/10/96
31	lubber	8	55.5	22	dustin	7	45.0	58	103	11/12/96
58	rusty	10	35.0	31	lubber	8	55.5	22	101	10/10/96
				31	lubber	8	55.5	58	103	11/12/96
				58	rusty	10	35.0	22	101	10/10/96
				58	rusty	10	35.0	58	103	11/12/96

R1		
sid	bid	day
22	101	10/10/96
58	103	11/12/96

8

Basic SQL Query

SELECT [DISTINCT] target-list

FROM relation-list

WHERE qualification

- **Relation-list:** A list of relation names (possibly with **range-variable** after each name).
- **Target-list:** A list of attributes of relations in relation-list
- **Qualification:** conditions on attributes (<, >, =, and, or, not, etc.)
- **DISTINCT:** optional keyword for duplicate removal.
 - Default = no duplicate removal!

9

How to evaluate a query?

SELECT [DISTINCT] target-list

FROM relation-list

WHERE qualification

- **Conceptual query evaluation** using relational operators:
 - 1) Compute the cross-product of relation-list.
 - 2) Discard resulting tuples if they fail qualifications.
 - 3) Delete attributes that are not in target-list. (called column-list)
 - 4) If DISTINCT is specified, eliminate duplicate rows.
- Only conceptual because of **inefficiency** computation
 - An optimizer can find better strategy

SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=103

10

Example of Conceptual Evaluation (1)

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=103
```

(1) Compute the cross-product of relation-list.

Sailors

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

X

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

11

Example of Conceptual Evaluation (2)

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=103
```

(2) Discard tuples if they fail qualifications.

Sailors X Reserves

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

12

Example of Conceptual Evaluation (3)

```
SELECT S.sname
FROM Sailors S, Reserves R
WHERE S.sid=R.sid AND R.bid=103
```

(3) Delete attribute columns that not in target-list.

Sailors X Reserves

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

A Note on Range Variables

```
SELECT S.sname
FROM Sailors as S, Reserves R
WHERE S.sid=R.sid AND bid=103
```

OR

```
SELECT sname
FROM Sailors, Reserves
WHERE Sailors.sid=Reserves.sid AND bid=103
```

- Really needed **range variables** only if the same relation appears twice in the FROM clause.

```
SELECT sname
FROM Sailors S, Reserves R1,
Reserves R2
WHERE S.sid = R1.sid AND
S.sid = R2.sid AND
R1.bid <> R2.bid
```

Find the sids of sailors who've reserved at least one boat

```
SELECT S.sid
FROM Sailors S, Reserves R
WHERE S.sid=R.sid
```

Sailors X Reserves

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

DISTINCT

- Find the names and ages of all sailors

```
SELECT S.sname, S.age
FROM Sailors S
```
- Add **DISTINCT** to this query?
- Replace S.sname by S.sid in the *SELECT* clause?
- Add *DISTINCT* to the above?

Sid	Sname	Rating	Age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Find sailors whose names begin and end with B and contain at least three characters.

```
SELECT S.age,
       age1=S.age-5,
       2*S.age AS age2
FROM   Sailors S
WHERE  S.sname LIKE 'B_%B'
```

- **AS** and **=** are two ways to name fields in result.
- **LIKE** for string matching.
 - **'_'** for one character
 - **'%'** for 0 or more characters.

Sid	Sname	Rating	Age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	20

Age	Age1	Age2
20	15	40

17

Find sid's of sailors who've reserved a red **or** a green boats.

```
SELECT DISTINCT S.sid
FROM   Sailors S, Boats B, Reserves R
WHERE  S.sid=R.sid AND R.bid=B.bid
       AND (B.color='red' OR B.color='green')
```

- **UNION**: work on two **union-compatible** sets of tuples

```
SELECT S.sid
FROM   Sailors S, Boats B, Reserves R
WHERE  S.sid=R.sid AND R.bid=B.bid AND B.color='red'
UNION
SELECT S.sid
FROM   Sailors S, Boats B, Reserves R
WHERE  S.sid=R.sid AND R.bid=B.bid AND B.color='green'
```

- Replace **OR** by **AND** in the first version?

18

Find sid's of sailors who've reserved a red and a green boat

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
INTERSECT
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='green'
```

- What do we get if we replace *INTERSECT* by *EXCEPT*?
 - (A Except B) returns tuples in A but not in B.
 - Find sids of all sailors who have reserved a red boat **but not** a green boat.

19

SET Construct: *UNION ALL*

- *UNION*, *INTERSECT*, and *EXCEPT* delete **duplicate by default**. (why?)
- To retain duplicates, use *UNION ALL*, *INTERSECT ALL*, or *EXCEPT ALL*.

Sid	Sname
71	Zorba
74	Horatio
74	Horatio
95	Bob

INTERSECT ALL

Sid	Sname
22	Dustin
71	Zorba
74	Horatio
74	Horatio

=

Sid	Sname
71	Zorba
74	Horatio
74	Horatio

20

Prior Knowledge: A Predictive Database For Developers

- <http://techcrunch.com/2012/09/11/prior-knowledge-a-predictive-database-for-developers/>

Step 1: Understand the data

Take a quick look at the contents of the `bank-data.csv` file to familiarize yourself with the dataset.

- 600 rows of customer data
- 7 demographic columns
- One column per financial product available for purchase
 - Savings account (`save_act`)
 - Checking account (`current_act`)
 - Mortgage (`mortgage`)
 - Personal equity plan (`pep`)
- Many missing values — not all values are known for all customers



age	sex	region	income	married	children	car	save_act	current_act	mortgage	pep
22	Male	Rural	8877.07	False	-	False	False	True	False	True
37	Female	Suburban	25304.3	True	2	True	False	False	-	False
66	Female	-	-	-	1	True	True	True	True	True
19	Male	Rural	10953.0	True	3	True	True	True	-	False
-	Male	Inner City	-	-	-	True	True	True	True	False

Nested Queries

- WHERE clause can itself contain an SQL **subquery**. (Actually, so can FROM and HAVING clauses.)
- Find names of sailors who've reserved boat #103:

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN (SELECT R.sid
                FROM Reserves R
                WHERE R.bid=103)
```

Subquery: finds sids who have reserved bid 103
- (`x IN B`) returns true when x is in set B.
 - To find sailors who've not reserved #103, use `NOT IN`.
- Nested loops Evaluation
 - For each Sailors tuple, check the qualification by computing the subquery.
 - Does the subquery result change for each new Sailor row?

Nested Queries with Correlation

```
SELECT S.sname
FROM   Sailors S
WHERE  EXISTS (SELECT *
               FROM   Reserves R
               WHERE  R.bid=103 AND S.sid=R.sid )
```

Correlation: subquery finds all reservations for bid 103 from current sid

- **EXISTS** is another set comparison operator, like **IN**.
 - (EXISTS S) returns true when S is not empty.
- What is the above query in English?
 - Find sailors who have reserved boat #103
- In case of correlation, subquery must be re-computed for each Sailors tuple.

23

Nested Queries with **UNIQUE**

Sailors(sid: integer, sname: string, rating: integer, age: real)
Boats(bid: integer, bname: string, color: string)
Reserves(sid: integer, bid: integer, day: date)

- (UNIQUE S) returns true if S has no duplicate tuples or S is empty.

```
SELECT S.sname
FROM   Sailors S
WHERE  UNIQUE (SELECT R.bid
               FROM   Reserves R
               WHERE  R.bid=103 AND S.sid=R.sid)
```

- What is the above query in English?
 - Finds sailors with at most one reservation for boat #103.
- Replace R.bid with *?
 - Finds sailors with at most one reservation for boat #103 in a given day.
 - (Simplify -> find all sailors)

24

More on Set-Comparison Operators

- Have seen *IN*, *EXISTS* and *UNIQUE*. Can also use *NOT IN*, *NOT EXISTS*, and *NOT UNIQUE*.
 - Also available: *op ANY*, *op ALL*, where *op* can be *>*, *<*, *=*, *≠*, *≤*, *≥*
 - (*a > ANY B*) returns true when *a* is greater than any one element in set *B*.
 - (*a > ALL B*) returns true when *a* is greater than all elements in set *B*.
- ```
SELECT *
FROM Sailors S
WHERE S.rating > ANY (SELECT S2.rating
 FROM Sailors S2
 WHERE S2.sname='Horatio')
```
- What is the above query in English?
    - Find sailors whose rating is greater than that of some sailor called Horatio.
  - What is the above query in English if *> ANY* is replaced by *> ALL*?
    - Find sailors whose rating is greater than all sailors called Horatio.

25

## Find sid's of sailors who've reserved a red and a green boat

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
INTERSECT
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='green'
```

- Rewrite *INTERSECT* with *IN*.
  - Strategy?

26

# Rewriting *INTERSECT* Using *IN*

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
 AND S.sid IN (SELECT S2.sid
 FROM Sailors S2, Boats B2, Reserves R2
 WHERE S2.sid=R2.sid AND R2.bid=B2.bid
 AND B2.color='green')
```

*Find sids who've reserved  
a green boat*

- Find *sid*'s of Sailors who've reserved red but not green boats (*EXCEPT*)
  - Replace *IN* with *NOT IN*.

27

## Aggregate Operators

- *COUNT* (\*)
- *COUNT* ( [*DISTINCT*] A)
  - A is a column
- *SUM* ( [*DISTINCT*] A)
- *AVG* ( [*DISTINCT*] A)
- *MAX* (A)
- *MIN* (A)
- Count the number of sailors

```
SELECT COUNT (*)
FROM Sailors S
```

28

*Find the average age of sailors with  
rating = 10*

*Sailors(sid: integer, sname: string, rating: integer,  
age: real)*

```
SELECT AVG (S.age)
FROM Sailors S
WHERE S.rating=10
```

29

*Count the number of different sailor  
names*

*Sailors(sid: integer, sname: string, rating: integer,  
age: real)*

```
SELECT COUNT (DISTINCT S.sname)
FROM Sailors S
```

30

## *Find the age of the oldest sailor*

*Sailors*(sid: integer, sname: string, rating: integer,  
age: real)

```
SELECT MAX(S.AGE)
FROM Sailors S
```

31

## *Find name and age of the oldest sailor(s)*

```
SELECT S.sname, MAX (S.age)
FROM Sailors S
```

- This is illegal, but why?
  - Cannot combine a column with a value (unless we use *GROUP BY*)

```
SELECT S.sname, S.age
FROM Sailors S
```

```
WHERE S.age = (SELECT MAX (S2.age) FROM Sailors S2)
```

- Okay, but not supported in every system
  - Convert a table (of a single aggregate value) into a single value for comparison

32



# GROUP BY and HAVING

- So far, aggregate operators are applied to all (qualifying) tuples.
  - Can we apply them to each of several groups of tuples?
- Example: find the age of the youngest sailor for each rating level.
  - In general, we don't know how many rating levels exist, and what the rating values for these levels are!
  - Suppose we know that rating values go from 1 to 10; we can write 10 queries that look like this:

For  $i = 1, 2, \dots, 10$ :

```
SELECT MIN (S.age)
FROM Sailors S
WHERE S.rating = i
```

33

## Find the age of the youngest sailor for each rating level

```
SELECT S.rating, MIN (S.age) as age
FROM Sailors S
GROUP BY S.rating
```

- (1) The sailors tuples are put into "same rating" groups.
- (2) Compute the Minimum age for each rating group.

| Sid | Sname  | Rating | Age  |
|-----|--------|--------|------|
| 22  | Dustin | 7      | 45.0 |
| 31  | Lubber | 8      | 55.5 |
| 85  | Art    | 3      | 25.5 |
| 32  | Andy   | 8      | 25.5 |
| 95  | Bob    | 3      | 63.5 |

| Rating | Age  |
|--------|------|
| 3      | 25.5 |
| 7      | 45.0 |
| 8      | 25.5 |

(2)

| Rating | Age  |
|--------|------|
| 3      | 25.5 |
| 3      | 63.5 |
| 7      | 45.0 |
| 8      | 55.5 |
| 8      | 25.5 |

(1)

Find the age of the youngest sailor for each rating level that has at least 2 members

```
SELECT S.rating, MIN (S.age) as age
FROM Sailors S
GROUP BY S.rating
HAVING COUNT(*) > 1
```

1. The sailors tuples are put into "same rating" groups.
2. Eliminate groups that have < 2 members.
3. Compute the Minimum age for each rating group.

| Sid | Sname  | Rating | Age  |
|-----|--------|--------|------|
| 22  | Dustin | 7      | 45.0 |
| 31  | Lubber | 8      | 55.5 |
| 85  | Art    | 3      | 25.5 |
| 32  | Andy   | 8      | 25.5 |
| 95  | Bob    | 3      | 63.5 |

| Rating | Age  |
|--------|------|
| 3      | 25.5 |
| 3      | 63.5 |
| 7      | 45.0 |
| 8      | 55.5 |
| 8      | 25.5 |

## Queries With GROUP BY and HAVING

```
SELECT [DISTINCT] target-list
FROM relation-list
WHERE qualification
GROUP BY grouping-list
HAVING group-qualification
```

```
SELECT S.rating, MIN (S.age) as age
FROM Sailors S
GROUP BY S.rating
HAVING S.rating > 5
```

- The *target-list* contains (i) attribute names (ii) terms with aggregate operations (e.g., AVG (S.age)).
- The attribute list (e.g., S.rating) in *target-list* must be in *grouping-list*.
- The attributes in *group-qualification* must be in *grouping-list*.

## Say if Attribute list is not in grouping-list

```
SELECT S.sname, S.rating, AVG
 (S.age) as age
FROM Sailors S
GROUP BY S.rating
HAVING COUNT(*) > 1
```

| Sid | Sname  | Rating | Age  |
|-----|--------|--------|------|
| 22  | Dustin | 7      | 45.0 |
| 31  | Lubber | 8      | 55.5 |
| 85  | Art    | 3      | 25.5 |
| 32  | Andy   | 8      | 25.5 |
| 95  | Bob    | 3      | 63.5 |

| Sname | Rating | Age  |
|-------|--------|------|
| ?     | 3      | 44.5 |
| ?     | 8      | 40.5 |

| Sname  | Rating | Age  |
|--------|--------|------|
| Art    | 3      | 25.5 |
| Bob    | 3      | 63.5 |
| Dustin | 7      | 45.0 |
| Lubber | 8      | 55.5 |
| Andy   | 8      | 25.5 |

37

## Say if Group qualification is not in grouping-list

```
SELECT S.rating, AVG (S.age) as age
FROM Sailors S
GROUP BY S.rating
HAVING S.sname ≠ 'Dustin'
```

Not in group-list

| Sid | Sname  | Rating | Age  |
|-----|--------|--------|------|
| 22  | Dustin | 7      | 45.0 |
| 31  | Lubber | 8      | 55.5 |
| 85  | Art    | 3      | 25.5 |
| 32  | Andy   | 8      | 25.5 |
| 95  | Bob    | 3      | 63.5 |

| Rating | Age |
|--------|-----|
|        |     |
|        |     |

| Sname  | Rating | Age  |
|--------|--------|------|
| Art    | 3      | 25.5 |
| Bob    | 3      | 63.5 |
| Dustin | 7      | 45.0 |
| Lubber | 8      | 55.5 |
| Andy   | 8      | 25.5 |

38

# Conceptual Evaluation

- Without *GROUP BY* and *HAVING*:
  - Compute cross-product of *relation-list*
  - Remove tuples that fail *qualification*
  - Delete unnecessary columns
- With *GROUP BY* and *HAVING*, continue with
  - Partition remaining tuples into groups by the value of attributes in *grouping-list* (specified in *GROUP-BY* clause)
  - Remove groups that fail *group-qualification* (specified in *HAVING* clause).
  - Compute one answer tuple per qualifying group.

39

*For each red boat, find the number of reservations for this boat*

```
SELECT B.bid, COUNT (*) AS
 num_reservations
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND
 B.color='red'
GROUP BY B.bid
```

```
SELECT B.bid, COUNT (*) AS
 num_reservations
FROM Boats B, Reserves R
WHERE R.bid=B.bid
GROUP BY B.bid
HAVING B.color='red'
```

- Illegal, why?
  - B.color does not appear in group-list

40

Find the age of the youngest sailor with age > 18 for each rating with *at least 2 sailors (of any age)*

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age > 18
GROUP BY S.rating
HAVING COUNT(S) > 1
```

- What is wrong?
  - COUNT(S) is counting tuples after the qualification (S.age > 18).
  - Eliminate groups, which have multiple sailors but only one sailor with age > 18.

- How to fix it?
  - Use subquery in the HAVING clause.

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age > 18
GROUP BY S.rating
HAVING
 1 < ANY (SELECT COUNT (*)
 FROM Sailors S2
 WHERE
 S.rating=S2.rating)
```

41

## Table Constraints

- Specify constraints over a single table
  - Useful when more general ICs than keys are involved.

```
CREATE TABLE Sailors
(sid INTEGER,
 sname CHAR(10),
 rating INTEGER,
 age REAL,
 PRIMARY KEY (sid),
 CHECK (rating >= 1
 AND rating <= 10)
```

- Constraints can be named.

```
CREATE TABLE Reserves
(sname CHAR(10),
 bid INTEGER,
 day DATE,
 PRIMARY KEY (bid, day),
 CONSTRAINT noInterlakeRes
 CHECK (`Interlake' ≠
 (SELECT B.bname
 FROM Boats B
 WHERE B.bid=bid)))
```

The boat 'Interlake' cannot be reserved

42

# Assertions: Constraints Over Multiple Tables

- Awkward and wrong!
  - If Sailors is empty, the number of Boats tuples can be anything!
- ASSERTION is the right solution; not associated with either table.

```
CREATE TABLE Sailors
(sid INTEGER,
 sname CHAR(10),
 rating INTEGER,
 age REAL,
 PRIMARY KEY (sid),
 CHECK
 ((SELECT COUNT (S.sid) FROM Sailors S)
 + (SELECT COUNT (B.bid) FROM Boats B) < 100)
```

*Number of boats  
plus number of  
sailors is < 100*

```
CREATE ASSERTION smallClub
CHECK
((SELECT COUNT (S.sid) FROM Sailors S)
 + (SELECT COUNT (B.bid) FROM Boats B) < 100)
```

43

# Triggers

- Trigger: procedure that starts automatically if specified changes occur to the DBMS
- A trigger has three parts:
  - Event (activates the trigger)
  - Condition (tests whether the triggers should run)
  - Action (what happens if the trigger runs)

```
CREATE TRIGGER incr_count AFTER INSERT ON Students // Event
WHEN (new.age < 18) // Condition
FOR EACH ROW
BEGIN // ACTION: a procedure in Oracle's PL/SQL syntax
 count := count + 1
END
```

44

# References

- Deutsch, Donald R., "The SQL Standard: How it Happened," Annals of the History of Computing, IEEE , vol.35, no.2, pp.72,75, April-June 2013

45

# Starwar Exercises

```
char(name, race, homeworld, affiliation)
 planets(name, type, affiliation)
timetable(cname, pname, movie, arrival, departure)
```

Which planet does Princess Leia go to in movie3?

```
SELECT distinct pname
FROM timetable
WHERE cname ='Princess Leia' and movie=3
```

46

# Starwar Exercises

char(name, race, homeworld, affiliation)  
planets(name, type, affiliation)  
timetable(cname, pname, movie, arrival, departure)

- How many people stay on Dagobah in movie 3?

```
SELECT count(*)
FROM timetable, characters
WHERE movie=3 and pname ='Dagobah' and
timetable.cname=characters.name and
characters.race='Human'
```

47

# Starwar Exercises

char(name, race, homeworld, affiliation)  
planets(name, type, affiliation)  
timetable(cname, pname, movie, arrival, departure)

- Who has been to his/her homeworld in movie 2?

```
SELECT distinct c.name
FROM characters c, timetable t
WHERE c.name=t.cname and t.pname=c.homeworld and
movie=2
```

48



# Starwar Exercises

char(name, race, homeworld, affiliation)

planets(name, type, affiliation)

timetable(cname, pname, movie, arrival, departure)

- Find all characters that have been on all planets of rebels.

```
SELECT name
```

```
FROM characters c
```

```
WHERE not exists (
```

```
 SELECT p.name FROM planets p
```

```
 WHERE affiliation='rebels' and p.name NOT IN
```

```
 (SELECT pname from timetable t where
 t.cname=c.name and t.pname=p.name))
```

49

# Starwar Exercises

char(name, race, homeworld, affiliation)

planets(name, type, affiliation)

timetable(cname, pname, movie, arrival, departure)

- Find distinct names of the planets visited by those of race "droid".

```
SELECT distinct t.pname
```

```
FROM characters c, timetable t
```

```
WHERE c.name=t.cname and c.race='droid'
```

50

# Starwar Exercises

char(name, race, homeworld, affiliation)

planets(name, type, affiliation)

timetable(cname, pname, movie, arrival, departure)

- For each character and for each neutral planet, how much time total did the character spend on the planet?

```
SELECT c.name, p.name, SUM(t.departure-t.arrival+1) as amount
```

```
FROM characters c, timetable t, planets p
```

```
WHERE t.cname=c.name and t.pname=p.name and p.affiliation='neutral'
```

```
GROUP BY c.name, p.name
```

51

## Division in SQL

- Find sailors who've reserved all boats.

- Strategy?

- Find all boats that have been reserved by a sailor
- Compare with all boats
- Do the sailor's reserved boats include all boats?
  - Yes → include this sailor
  - No → exclude this sailor

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
 ((SELECT B.bid
 FROM Boats B)
 EXCEPT
 (SELECT R.bid
 FROM Reserves R
 WHERE R.sid=S.sid))
```

(A EXCEPT B) returns  
tuples in A but not in  
B.

52

# Division in SQL

- Can you do it the hard way, without *EXCEPT* & with *NOT EXISTS*?
- Strategy:
  - For each sailor, check that there is no boat that has not been reserved by this sailor.

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS (SELECT B.bid
 FROM Boats B
 WHERE NOT EXISTS (SELECT
R.bid
 FROM Reserves R
 WHERE R.bid = B.bid AND R.sid = S.sid)
```

Sailors

| <u>sid</u> | sname  | rating | age  |
|------------|--------|--------|------|
| 22         | dustin | 7      | 45.0 |
| 31         | lubber | 8      | 55.5 |

Boats

| <u>bid</u> | bname | color |
|------------|-------|-------|
| 101        | xyz   | red   |
| 103        | abc   | green |

Reserves

| <u>sid</u> | <u>bid</u> | day      |
|------------|------------|----------|
| 22         | 101        | 10/10/96 |
| 31         | 101        | 11/12/96 |
| 31         | 103        | 12/12/96 |